

Intro to R: Part 2

IBIO 851

Sept 8 2016

Suggested
reading:

**Chapter 2 in
Biostatistical
Design &
Analysis
Using R
(Logan)**

Goals for today

- Review key basics
- Importing data
- Viewing data
- Basics of data manipulation
- Reproducibility in science
- Intro to R Markdown

Data types

- R has a wide variety of data types
 - Scalars, vectors, matrices, dataframes, lists
- Vectors
 - a <- c(1,2,5.3,6,-2,4) # numeric
 - b <- c("one","two","three") # character
 - d <- c(TRUE,TRUE,TRUE,FALSE,TRUE) #logical
- Refer to elements of a vector using subscripts
 - a[2:4] # 2nd through 4th elements of vector

Data types

- **Matrices**: all columns in a matrix must have the same mode (numeric, character, etc.) and same length

```
mymatrix <- matrix(vector, nrow=r, ncol=c,  
byrow=FALSE,dimnames=list(char_vector_rownames,  
char_vector_colnames))
```

- **byrow**=TRUE indicates that the matrix should be filled by rows
- **byrow**=FALSE indicates that the matrix should be filled by columns (the default)
- **dimnames** provides optional labels for the columns and rows

Data types

- Matrices
- Generates 5 x 4 numeric matrix:
`y<-matrix(1:20, nrow=5,ncol=4)`
- Identify rows, columns or elements using subscripts.
 - `x[,4]` # 4th column of matrix
 - `x[3,]` # 3rd row of matrix
 - `x[2:4,1:3]` # rows 2,3,4 of columns
1,2,3

Data types

- **Arrays**: similar to matrices, but can have more than 2 dimensions
 - See `help(array)` for details
- **Data frames**: more general than a matrix in that it can include non-numbers

```
> d <- c(1,2,3,4)
> e <- c("red", "white", "red", NA)
> f <- c(TRUE,TRUE,TRUE,FALSE)
> mydata <- data.frame(d,e,f)
> names(mydata) <- c("ID","Color","Passed") #variable names
>
> mydata
  ID Color Passed
1  1   red   TRUE
2  2 white   TRUE
3  3   red   TRUE
4  4  <NA> FALSE
>
```

Data types

- Dataframes

- There are a variety of ways to identify elements of a dataframe

`myframe[3:5]` # columns 3,4,5 of dataframe

`myframe[c("ID","Age")]` # columns ID & Age

`myframe$X1` # variable x1

Data types

- **Lists**: An ordered collection of objects (components)
- A list allows you to gather a variety of (possibly unrelated) objects under one name
- Example of a list with 4 components:

```
> w <- list(name="Fred", mynumbers=d, mymatrix=y, age=5.3)
>
> w
$name
[1] "Fred"

$mynumbers
[1] 1 2 3 4

$mymatrix
      [,1] [,2] [,3] [,4]
[1,]    1    6   11   16
[2,]    2    7   12   17
[3,]    3    8   13   18
[4,]    4    9   14   19
[5,]    5   10   15   20

$age
[1] 5.3

> w[[2]]
[1] 1 2 3 4
>
```


Useful functions

- `length(object)` #number of elements
- `str(object)` #structure of an object
- `names(object)` #names (e.g. variables)
- `cbind(object, object)` #combine objects
as columns
- `rbind(object, object)` #combine objects
as rows
- `ls()` #list current objects stored in memory

Importing data

- From a **comma delimited text file**:
 - First row contains variable names, comma is separator
- (1) `mydata <- read.table("c:/mydata.txt", header=TRUE, sep="," , row.names="id")`
- (2) `mydata <- read.csv("c:/mydata.csv", header=TRUE, sep="," , row.names="id")`
- **header=TRUE** means that you have column headings in the first row
 - **sep=","** means that the data is comma-delimited
 - **row.names="id"** means that the number of each row is now the ID

Importing data

- From **Excel**:
 - The best way to read an Excel file is to export it to a comma delimited file and import it using the previous method

Importing data

- **Manually** in R:
 - Create a dataframe from scratch
 - Only possible with very small datasets

```
> age <- c(25, 30, 56)
> gender <- c("male", "female", "male")
> weight <- c(160, 110, 220)
> mydata <- data.frame(age,gender,weight)
> mydata
```

	age	gender	weight
1	25	male	160
2	30	female	110
3	56	male	220

Viewing data

- There are a number of functions for listing the contents of an object or dataset
- Here are commonly used ones (some you've seen already):

```
# list objects in the working  
environment  
ls()
```

```
# list the variables in mydata  
names(mydata)
```

```
# list the structure of mydata  
str(mydata)
```

```
# list levels of factor v1 in mydata  
levels(mydata$v1)
```

```
# dimensions of an object  
dim(object)
```

Viewing data

```
# class of an object (numeric, matrix,  
  dataframe, etc)  
class(object)  
  
# print mydata (this shows up in the  
  console)  
mydata  
  
# print first 10 rows of mydata  
head(mydata, n=10)  
  
# print last 5 rows of mydata  
tail(mydata, n=5)
```

Data manipulation

- Once you have imported your data, you will want to manipulate it into a useful form
- This includes creating new variables, **sorting & merging** datasets, **aggregating** data, **reshaping** data, and **subsetting** datasets

Data manipulation

- Each of these activities usually involves the use of R's built-in functions & operators
- Sometimes you have to write your own function!
- Other times you will need to convert variables or datasets to another type (e.g. numeric to character or matrix to dataframe)

Creating new variables

- Use the **assignment operator** `<-` to create new variables

```
mydata$sum <- mydata$x1 + mydata$x2  
mydata$mean <- (mydata$x1 + mydata$x2)/2
```

- Sometimes you will need to **recode** variables (i.e. create another variable from input data)

```
> age<-c(12, 16, 13, 32, 46, 78, 45, 98, 45)  
> age[age>75]<- "Elder"  
> age  
[1] "12"    "16"    "13"    "32"    "46"    "Elder" "45"    "Elder" "45"  
> age[age>45 & age <=75]<- "Middle Aged"  
> age  
[1] "12"    "16"    "13"    "32"    "Middle Aged" "Elder"  "45"    "Elder"  "45"  
> age[age<=45]<- "Young"  
> age  
[1] "Young"  "Young"  "Young"  "Young"  "Middle Aged" "Elder"  "Young"  "Elder"  "Young"
```

Merging

- To merge two dataframes (datasets) horizontally, use the `merge` function
- In most cases, you join two dataframes by one or more common key variables (i.e., an inner join)

```
# merge two dataframes by ID
```

```
total <- merge(dataframeA, dataframeB, by="ID")
```

```
# merge two dataframes by ID & Country
```

```
total <- merge(dataframeA, dataframeB,  
by=c("ID", "Country"))
```

Practicing in R

- In class lab 2 & vector practice

Reproducibility in science

Reproducibility

- **Reproducibility** is the ability of an entire experiment (or statistical analysis) to be duplicated, either by the same researcher or by someone else working independently
- It is one of the main principles of the scientific method
- It has come to the forefront recently



Explore ▾

LATEST

MOST VIEWED

NEWS
Debate accelerates on universe's expansion speed
BY EMILY CONOVER JULY 22, 2016

NEWS
How dinosaurs hopped across an ocean
BY THOMAS SUMNER JULY 22, 2016

SCIENCE STATS
U.S. lags in robotics
BY ALEX MADDON

NEWS
Yeasts hide in partnerships
BY SUSAN MILIUS

SCIENCE TICKER
Getting rid of stopping snail
BY AMY MODERNOTT

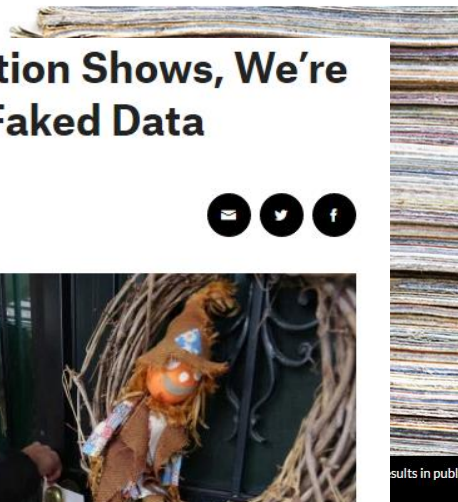
SCIENCE NEWS FOR STUDENTS
The shocking

NEWS
Humans, birds collaborate
BY BRUCE BOWER

FEATURE SCIENCE & SOCIETY, NUMBERS, 2015 TOP 25

Year in review: Scientists tackle the irreproducibility problem

Banking on experimental results that don't hold up to replication is expensive
BY TINA HESMAN SAEY 7:00AM, DECEMBER 15, 2015



As A Major Retraction Shows, We're All Vulnerable To Faked Data

By Carl Bielik

Filed under Data



A Colorado political canvasser in 2014. A major study showing the power of personal contact by campaigns was retracted on Tuesday. BRENNAN LINSLEY/AP

A political scientist on Tuesday said he was retracting a paper he'd co-authored — one with wide influence on how campaigns can change public opinion — when faced with evidence that the paper's central finding was based on polling that probably never happened.

nature


International weekly journal of science

Home | News & Comment | Research | Careers & Jobs | Current Issue | Archive | Audio & Video | For Authors

Archive > Specials and supplements archive > Challenges in irreproducible research

SPECIAL

See all specials



CHALLENGES IN IRREPRODUCIBLE RESEARCH

Science moves forward by corroboration — when researchers verify others' results. Science advances faster when people waste less time pursuing false leads. No research paper can ever be considered to be the final word, but there are too many that do not stand up to further study.

There is growing alarm about results that cannot be reproduced. Explanations include increased levels of scrutiny, complexity of experiments and statistics, and pressures on researchers. Journals, scientists, institutions and funders all have a part in tackling reproducibility. *Nature* has taken substantive steps to improve the transparency and robustness in what we publish, and to promote awareness within the scientific community. We hope that the articles contained in this collection will help.

▼ Editorial ▼ Features ▼ News and analysis ▼ Comment

▼ Perspectives and reviews

As A Major Retraction Shows, We're All Vulnerable To Faked Data

By [Carl Bielik](#)

Filed under [Data](#)



A Colorado political canvasser in 2014. A major study showing the power of personal contact by campaigns was retracted on Tuesday. BRENNAN LINSLEY/AP

A political scientist on Tuesday said he was retracting a paper he'd co-authored — one with wide influence on how campaigns can change public opinion — when faced with evidence that the paper's central finding was based on polling that probably never happened.

The [article](#), published last December in Science Magazine by UCLA graduate student [Michael J. LaCour](#) and Columbia University political scientist [Donald P. Green](#), appeared to show that an in-person conversation with an openly gay person made voters feel much more positively about same-sex marriage, an effect that persisted and even spread to the people those voters lived with, who weren't part of the conversation. The result of that purported effect was an affirmation of the power of human contact to overcome disagreement.

By describing personal contact as a powerful political tool, the paper influenced many campaigns and activists to shift their approach to emphasize the power of the personal story. The study was featured by [Bloomberg](#), on "[This American Life](#)" and in [activists' playbooks](#), including those used by backers of an Irish [constitutional referendum](#) up for a vote Friday that would legalize same-sex marriage.

"How to convince anyone to change their mind on a divisive issue in just 22 minutes — with science," was [one catchy headline](#) on a [Business Insider story](#) about the study. (The article was updated Wednesday with news of the retraction.)

Now that the underlying data appears to be fallacious, and Green has asked to retract the study (in a letter to Science and in his online [CV](#)), the study reveals different lessons. It shows how easily a scientist can invent data to show a desired result. It also shows how other scientists looking to replicate the result, with access to the original data, can quickly expose bad research. It took the authors of a study debunking the Science paper just two days to write their findings after they first noticed anomalies in the research. Their

Noteworthy irreproducible results


- **Hideyo Noguchi** became famous for correctly identifying the bacterial agent of syphilis, but also claimed he could culture this agent in the lab
 - Nobody has been able to reproduce this result
- **MMR vaccine** controversy
 - A study in the Lancet (peer-reviewed med journal) claiming the MMR vaccine caused autism was revealed to be fraudulent

Going back to p-values

- A 2015 study wanted to assess the rate & predictors of reproducibility in the field of psychology (but this can apply to any field!)
- The authors looked at whether the following factors were predictive of the likelihood that a replication study would be statistically significant (given that the original study already obtained stat. sig. [SS] results):
 - Characteristics of investigator
 - Hypotheses
 - Analytical methods
 - P-value in original study

Going back to p-values

- Most factors did not predict reproducibility



P-values & reproducibility take-home messages

- The exact p-value matters—not just whether a result is significant or not
- A p-value near 0.05 isn't worth much by itself
- Replication is crucial

What does
this have to
do with us?

- To maintain the reproducibility of research, certain protocols should be followed like second nature
 - Maintaining well-annotated code
 - Noting final version of codes and models
 - Knowing which code produces the final figures that appear in your papers
 - Saving workspaces so that you can reproduce the same parameter estimates you report

What does
this have to
do with us?

- One method to achieve these goals is **R Markdown**
- **R Markdown** is a way to embed R code chunks in a shareable document, making it easy to create reproducible web-based reports

R Markdown Tour

What is R Markdown?



- Creating docs with R Markdown starts with a .Rmd file (not .R)
 - Contains a combination of R code chunks and simple text formatting
- .Rmd file fed to knitr, which executes all R code chunks and creates new .md doc including R code and output

What is R Markdown?



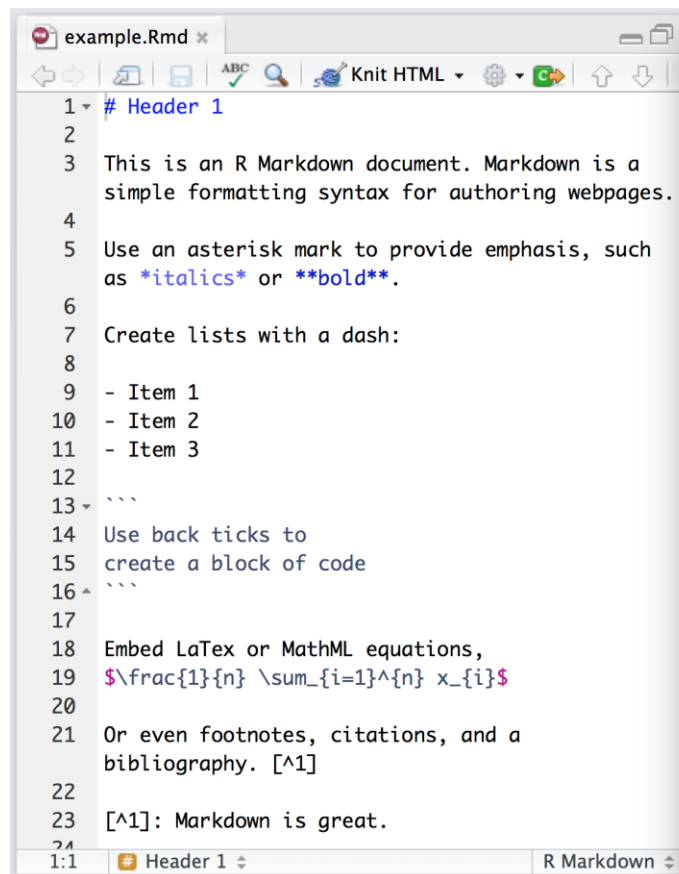
- The .md file is then processed by pandoc which creates a finished webpage, PDF, Word doc, etc.
- Rstudio includes a 'Knit' button that enables you to render a .Rmd file and preview it with a single click

R Markdown basics

- Installation
 - `install.packages("rmarkdown")`
- Simply go to File, New File, R Markdown to start a new .Rmd file

R Markdown basics

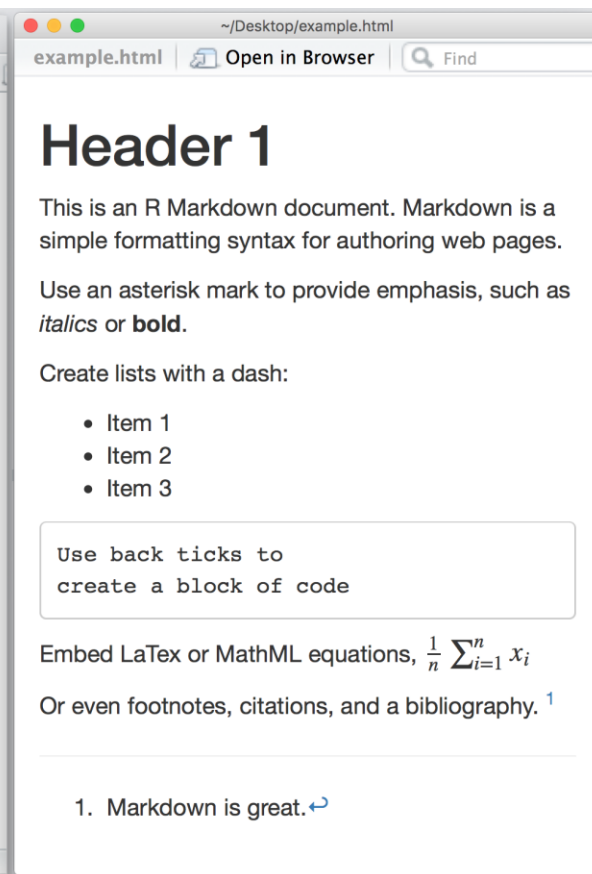
Before



```
1 # Header 1
2
3 This is an R Markdown document. Markdown is a
4 simple formatting syntax for authoring webpages.
5 Use an asterisk mark to provide emphasis, such
6 as italics or bold.
7 Create lists with a dash:
8
9 - Item 1
10 - Item 2
11 - Item 3
12
13 ```
14 Use back ticks to
15 create a block of code
16 ```
17
18 Embed LaTeX or MathML equations,
19 
$$\frac{1}{n} \sum_{i=1}^n x_i$$

20
21 Or even footnotes, citations, and a
22 bibliography. [^1]
23
24 [^1]: Markdown is great.
```

After



R Markdown basics

- R code chunks
 - Can be embedded within fenced code regions

The screenshot displays the RStudio interface with two windows. The left window, titled 'chunks.Rmd', shows the source R Markdown code. A teal circle highlights the fenced code chunk starting at line 6. The right window, titled 'RStudio: Preview HTML', shows the rendered output of the document.

chunks.Rmd Source Code:

```
1 R Code Chunks
2 =====
3
4 With R Markdown, you can insert R code
5 chunks including plots:
6 ```{r qplot, fig.width=4, fig.height=3,
7   message=FALSE}
8 # quick summary and plot
9 library(ggplot2)
10 summary(cars)
11 qplot(speed, dist, data=cars) +
12   geom_smooth()
13 ```
```

RStudio: Preview HTML Output:

R Code Chunks

With R Markdown, you can insert R code chunks including plots:

```
# quick summary and plot
library(ggplot2)
summary(cars)
```

##	speed	dist
##	Min. : 4.0	Min. : 2
##	1st Qu.:12.0	1st Qu.: 26
##	Median :15.0	Median : 36
##	Mean :15.4	Mean : 43
##	3rd Qu.:19.0	3rd Qu.: 56
##	Max. :25.0	Max. :120

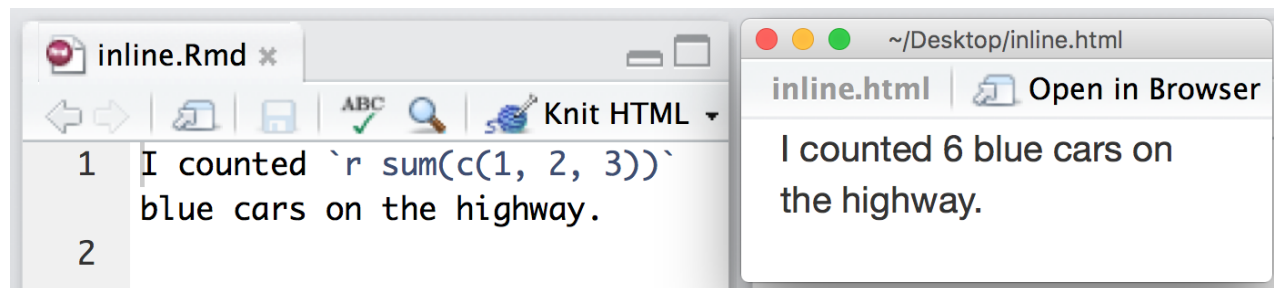
```
qplot(speed, dist, data = cars) + geom_smooth()
```

dist

speed

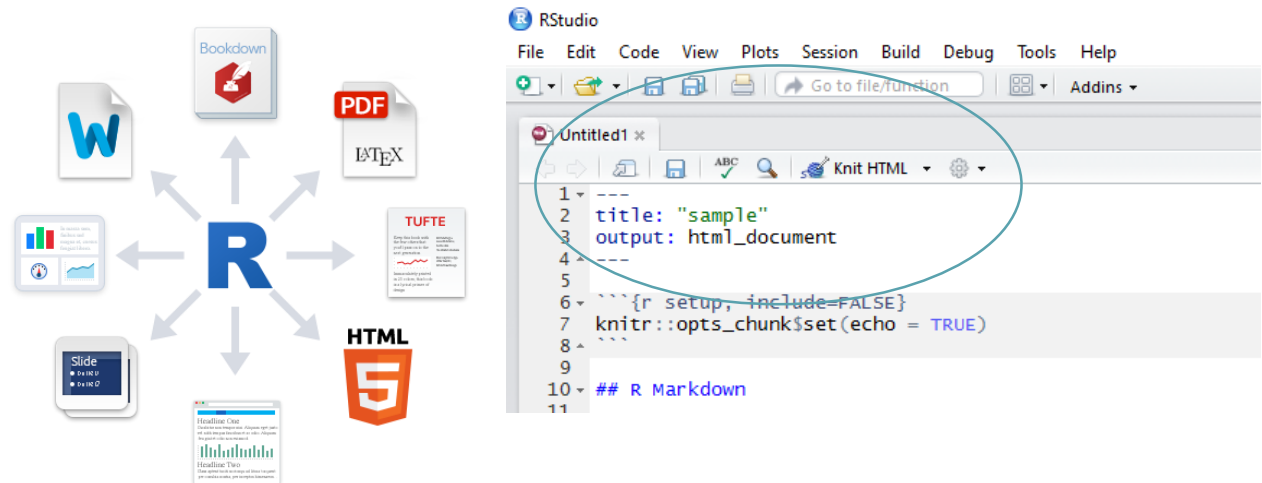
R Markdown basics

- Inline R code
 - You can evaluate R expressions inline by enclosing the expression within a single back-tick qualified with 'r'



R Markdown basics

- Rendering output
 - Use the Knit HTML button in upper left to finalize document
 - Specify what format you prefer when dialog box pops open under File > New
 - You can also specify format under output style at top of doc



R Markdown basics

- Let's open RStudio to try a .Rmd tutorial
- Assignment #1 due next Thurs (9/15)—can get started on it today